

Table of Contents

1. Part 1: Basic Concepts (Theory)
2. Part 2: Application and Analysis (Theory)
3. Part 3: Proof and Complexity (Theory)
4. Interval Scheduling: Implementation and Empirical Analysis

Repository: [GitHub](#)

Part 1: Basic Concepts

1. Three Primary Criteria for Algorithm Design

1. **Correctness:** Produces correct output for all valid inputs and terminates.
 - An algorithm must produce the **right answer** for every possible input within its domain.
 - Correctness can be formal (proved mathematically) or empirical (tested extensively).
 - Without correctness, an algorithm is fundamentally useless, regardless of speed or elegance.
 - Example: A sorting algorithm must produce a sorted array for all possible input arrays (any order, any values, any size).
2. **Efficiency:** Uses time and space resources effectively (Big-O time/space).
 - Time efficiency measures how the algorithm's runtime grows with input size n (e.g., $O(n \log n)$ is better than $O(n^2)$).
 - Space efficiency measures memory usage, critical for large-scale problems (e.g., in-place $O(1)$ is better than $O(n)$ auxiliary space).
 - Efficiency determines whether an algorithm is practical: a correct algorithm with $O(2^n)$ runtime is useless for $n > 30$.
 - Example: Sorting with $O(n \log n)$ (merge sort, quick sort) is preferred over $O(n^2)$ (bubble sort, insertion sort) for large datasets.
3. **Clarity/Simplicity:** Easy to understand, implement, and maintain.
 - Clear algorithms are less prone to bugs during implementation and easier to debug when issues arise.
 - Simple code is maintainable: future developers (or your future self) can understand and modify it without confusion.
 - There's a trade-off: sometimes the most efficient algorithm is less intuitive (e.g., fast Fourier transform vs. naive convolution).
 - In practice, choosing between a slightly slower but clear solution vs. a faster but complex one depends on the project's constraints.

- Example: A clear $O(n^2)$ solution may be preferable to an obscure $O(n \log n)$ solution if the data size is small and clarity matters more.

2. Five Basic Data Structure Types and Complexities

Data Structure	Avg Access	Avg Search	Avg Insert	Avg Delete	Avg Space
Array	$O(1)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$
Linked List	$O(n)$	$O(n)$	$O(1)^*$	$O(1)^*$	$O(n)$
Hash Table	$O(1)$	$O(1)$	$O(1)$	$O(1)$	$O(n)$
Binary Search Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(n)$
Heap	$O(n)$	$O(n)$	$O(\log n)$	$O(\log n)$	$O(n)$

*After locating the position to insert/delete

3. Objective of the Interval Scheduling Problem

Select a **maximum-size subset** of pairwise non-overlapping intervals. Formally:

- Input: $I = \{(s_i, f_i)\}$ with $s_i < f_i$
- Compatibility: Two intervals (s_i, f_i) and (s_j, f_j) are compatible if $f_i \leq s_j$ or $f_j \leq s_i$
- Goal: maximize $|S|$ where $S \subseteq I$ and all intervals in S are pairwise compatible

4. Three Greedy Strategies for Interval Scheduling

1. **Earliest Finish Time (EFT)**: sort by increasing finish time f_i (optimal).
2. **Earliest Start Time (EST)**: sort by increasing start time s_i (not always optimal).
3. **Shortest Duration (SD)**: sort by increasing $(f_i - s_i)$ (not always optimal).

5. Algorithm vs. Heuristic

Aspect	Algorithm	Heuristic
Correctness	Guaranteed for all inputs	Not guaranteed
Output	Always correct	Often good but can be suboptimal
Proof	Formal proof possible	Empirical validation only
Example	Dijkstra's shortest path	Nearest Neighbor for TSP

Part 2: Application and Analysis

1. Earliest Finish Time (EFT) Algorithm Steps

Algorithm Description: The Earliest Finish Time (EFT) greedy algorithm selects intervals greedily by always choosing the interval that finishes earliest among all compatible options. This ensures maximum room for future intervals.

Detailed Steps:

1. **Sort intervals by increasing finish time (f_i in ascending order).**
 - This preprocessing step is crucial: it orders intervals so we can greedily select from the earliest-finishing available intervals.
 - Sorting takes $O(n \log n)$ time, which dominates the algorithm's complexity.
2. **Initialize an empty selected set and $\text{lastfinish} = -\infty$.**
 - The selected set tracks which intervals we've chosen.
 - lastfinish tracks the finish time of the last selected interval (initialized to $-\infty$ so the first interval always qualifies).
3. **Scan in order and select each interval whose start is $\geq \text{lastfinish}$.**
 - For each interval in sorted order, check if it's compatible with previously selected intervals.
 - Compatibility is guaranteed because we track only lastfinish : if $s_i \geq \text{lastfinish}$, then interval i doesn't overlap with any selected interval.
 - This greedy choice—picking the earliest-finishing interval—maximizes the remaining time window for future intervals.
4. **Update lastfinish each time an interval is selected.**
 - After selecting interval i , update $\text{lastfinish} = f_i$ to ensure future selections are compatible.

Time Complexity: $O(n \log n)$ due to sorting.

2. Why Nearest Neighbor for TSP is a Heuristic

The Traveling Salesman Problem (TSP) asks: given n cities and distances between all pairs, find the shortest route visiting each city exactly once and returning to the start.

Nearest Neighbor (NN) Algorithm:

1. Start at an arbitrary city.
2. Repeatedly visit the unvisited city closest to the current city.
3. Return to the start city.

Why NN is a Heuristic (Not an Algorithm):

- **No Correctness Guarantee:** NN makes locally optimal choices (always picking the nearest unvisited city) that can lead to globally suboptimal tours.
- **No Optimality Bound:** Unlike approximation algorithms, NN has no guarantee.

- **NP-hard Problem:** TSP is NP-hard. NN is polynomial-time ($O(n^2)$), trading optimality for speed.

3. Explain the Concept of Loop Invariant and Its Relation to Mathematical Induction

A **loop invariant** is a property that holds before, during, and after each iteration of a loop. Proving the invariant uses the same structure as mathematical induction:

- **Base case:** invariant holds before the loop starts (at iteration 0).
- **Inductive step:** if it holds before iteration i , it holds after iteration i (before iteration $i + 1$).
- **Termination:** when the loop exits, the invariant implies the algorithm's correctness.

Connection to Induction: Loop invariants directly correspond to inductive proofs. For the EFT algorithm, the invariant might be: "After processing the first i intervals, the selected set contains the maximum-size compatible subset of those i intervals." We prove this by induction on i , and when $i = n$, the algorithm's correctness follows.

4. What Does It Mean for a Problem Like TSP to Be NP-hard?

A problem is **NP-hard** if it is at least as hard as the hardest problems in NP. For TSP, this means:

- **Exact solutions require exponential-time search** in the worst case.
- **No known polynomial-time exact algorithm** exists (and none likely will unless $P = NP$).
- **Verification is fast:** if someone gives you a proposed tour, you can check its length in polynomial time.
- **Consequence:** For large n , exact algorithms become impractical; heuristics and approximations are used instead.

5. How Does Modularity Contribute to Ease of Implementation?

Modularity decomposes an algorithm into independent, self-contained components, making it:

- **Easier to implement:** each module can be coded and tested separately.
- **Easier to debug:** if a bug exists, it's localized to a specific module.
- **Easier to optimize:** individual components can be optimized without affecting others.
- **Easier to reuse:** a module designed for one problem can often be used in another.

- **Easier to understand:** the overall logic is clearer when broken into simple pieces.

Part 3: Proof and Complexity

1. Proof by Contradiction: EFT Optimality

Claim: The Earliest Finish Time (EFT) greedy algorithm selects a maximum-size set of compatible intervals.

Proof by Contradiction:

Assume for contradiction that there exists an input where EFT selects fewer intervals than some optimal solution OPT. Let e_1, e_2, \dots, e_k be the intervals selected by EFT (in order of finish time), and o_1, o_2, \dots, o_m be the intervals in OPT, with $m > k$.

Since EFT is greedy, e_1 has the earliest finish time among all intervals. We can show by induction that for each $i \leq k$, the interval e_i finishes no later than o_i . This follows because:

- e_1 finishes earliest overall, so $f_{e_1} \leq f_{o_1}$.
- If $f_{e_i} \leq f_{o_i}$, then e_{i+1} (the next available interval after e_i) can start no earlier than e_i finishes, whereas o_{i+1} can start no earlier than o_i finishes. Since $f_{e_i} \leq f_{o_i}$, interval e_{i+1} finishes no later than o_{i+1} .

By induction, $f_{e_k} \leq f_{o_k}$. But then EFT could have selected o_{k+1} (which is compatible with e_1, \dots, e_k since it's compatible with o_1, \dots, o_k and $f_{e_k} \leq f_{o_k}$), contradicting that EFT selected only k intervals. Therefore, EFT must be optimal. \square

2. Counterexample: Shortest Duration (SD) Fails

Claim: The Shortest Duration (SD) greedy strategy—which sorts intervals by duration ($f_i - s_i$) and greedily selects the shortest—is NOT always optimal.

Counterexample:

- **A:** $[0, 10]$ (duration 10)
- **B:** $[1, 2]$ (duration 1)
- **C:** $[3, 11]$ (duration 8)

SD Algorithm Execution:

1. Sort by duration: B (1), C (8), A (10).
2. Select B: lastfinish = 2.
3. Consider C: $s_C = 3 \geq 2$, so select C: lastfinish = 11.
4. Consider A: $s_A = 0 < 11$, so REJECT A.

SD Result: 2 intervals (B, C).

Optimal Solution: Select B and A, or select A alone. Since B and A are compatible (B ends at 2, A starts at 0 but ends at 10, and B [1, 2] is after A starts), we can select both. Result: 2 intervals. Actually, EFT would select B and A, also 2 intervals.

Better Counterexample:

- **A:** [0, 5] (duration 5)
- **B:** [1, 2] (duration 1)
- **C:** [3, 4] (duration 1)
- **D:** [5, 6] (duration 1)

SD: Selects B (1), then C (compatible), then rejects D (overlaps with C), then rejects A. Result = 2 intervals (B, C).

Optimal (EFT): Selects B, C, D (all compatible). Result = 3 intervals.

Therefore, SD is not optimal. \square

3. Exhaustive Search Complexity for 20-point Robot Tour

For a Traveling Salesman Problem with 20 cities, the number of possible tours (treating clockwise and counterclockwise as equivalent) is:

$$(n - 1)!/2 = 19!/2 \approx 1.22 \times 10^{17}$$

Assuming a modern computer can evaluate 10^9 tours per second, this would require:

$$\frac{1.22 \times 10^{17}}{10^9} \approx 1.22 \times 10^8 \text{ seconds} \approx 3.9 \text{ years}$$

In practice, exhaustive search is far slower due to overhead, making it impractical for $n \geq 20$.

4. Why Induction Beats Trial-and-Error for Proving Correctness

Trial-and-Error (Testing): Can only validate the algorithm on a finite number of test cases. No matter how many inputs you test, you cannot cover all possible inputs.

Example: after testing an algorithm on 1 million inputs, you cannot guarantee it works on the 1-millionth-and-first input.

Mathematical Induction: Establishes correctness for **all valid inputs** by proving:

1. **Base case:** the algorithm is correct for the smallest input.
2. **Inductive step:** if the algorithm is correct for input of size k , then it's correct for size $k + 1$.

This proof covers infinitely many cases and is the only rigorous guarantee.

5. Complexity Trade-offs: Quicksort vs. Mergesort

Criterion	Quicksort	Mergesort
Average Time	$O(n \log n)$	$O(n \log n)$
Worst-case Time	$O(n^2)$	$O(n \log n)$
Space Complexity	$O(\log n)$ (in-place recursion)	$O(n)$ (extra arrays)
Stability	Not stable	Stable
Practical Speed	Faster on average (low constant factor)	Slower in practice (higher constant factor)

Trade-off Summary: Quicksort is faster on average and uses less memory, but can degrade to $O(n^2)$ on adversarial inputs. Mergesort guarantees $O(n \log n)$ worst-case time but requires $O(n)$ extra space. Choice depends on problem constraints (guaranteed vs. average performance, memory availability).

Interval Scheduling: Implementation and Empirical Analysis

This section implements all required algorithms, datasets, experiments, and plots for the interval scheduling assignment.

Objectives

1. Implement greedy algorithms (EFT, EST, SD)
2. Implement exhaustive optimal solver
3. Generate controlled synthetic datasets with varying overlap regimes
4. Measure empirical runtime and validate Big-O complexity
5. Analyze solution quality vs. optimal solutions
6. Visualize results and draw conclusions

Notebook Contents

This notebook integrates all required components:

- **Dataset Generation** (Section 2): Controlled synthetic interval generation with overlap regimes
- **Greedy Algorithms** (Section 3): EFT, EST, SD implementations with compatibility checking
- **Exhaustive Solver** (Section 4): Optimal subset enumeration algorithm for validation

- **Benchmarking Framework** (Section 5): Runtime measurement across input sizes with statistical analysis
- **Complexity Validation** (Section 6): Log-log and normalized runtime plots for Big-O analysis
- **Quality Analysis** (Section 7): Solution optimality comparison across overlap regimes

1. Import Required Libraries

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import time
from itertools import combinations
import seaborn as sns
from scipy import stats

# Set random seed for reproducibility
np.random.seed(42)

# Configure plotting style
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette("husl")

print("All libraries imported successfully!")
```

All libraries imported successfully!

2. Dataset Generation

Time Horizon Formula: $T = \alpha \cdot n \cdot D$

- α controls overlap density (0.1 = high, 1 = medium, 5 = low)
- n = number of intervals
- D = maximum interval duration

Dataset Generation Process: For each interval i :

- Start time: $s_i \sim \text{Uniform}[0, T]$
- Duration: $d_i \sim \text{Uniform}[1, D]$
- Finish time: $f_i = s_i + d_i$

```
In [2]: def generate_interval_dataset(n, alpha=1.0, D=10):
    """
    Generate a set of n random intervals with configurable overlap density
    """
    T = alpha * n * D
    intervals = []
    for _ in range(n):
        start = np.random.uniform(0, T)
        duration = np.random.uniform(1, D)
        finish = start + duration
        intervals.append((start, finish))
```

```

    return intervals

def get_overlap_density_label(alpha):
    """Return descriptive label for overlap density parameter alpha."""
    if alpha <= 0.15:
        return "High Overlap ( $\alpha \approx 0.1$ )"
    elif alpha <= 2:
        return "Medium Overlap ( $\alpha \approx 1$ )"
    else:
        return "Low Overlap ( $\alpha \approx 5$ )"

```

3. Greedy Algorithms Implementation

Greedy Algorithm Strategies

1. **Earliest Finish Time (EFT)** - optimal
2. **Earliest Start Time (EST)** - suboptimal in worst case
3. **Shortest Duration (SD)** - suboptimal in worst case

```

In [3]: def are_compatible(interval1, interval2):
    """Check if two intervals are compatible (non-overlapping)."""
    s1, f1 = interval1
    s2, f2 = interval2
    return f1 <= s2 or f2 <= s1

def earliest_finish_time(intervals):
    """Earliest Finish Time (EFT) Greedy Algorithm – OPTIMAL."""
    if not intervals:
        return 0, []
    sorted_intervals = sorted(enumerate(intervals), key=lambda x: x[1][1])
    selected = []
    last_finish_time = float('-inf')
    for idx, (start, finish) in sorted_intervals:
        if start >= last_finish_time:
            selected.append(idx)
            last_finish_time = finish
    return len(selected), sorted(selected)

def earliest_start_time(intervals):
    """Earliest Start Time (EST) Greedy Algorithm – Suboptimal."""
    if not intervals:
        return 0, []
    sorted_intervals = sorted(enumerate(intervals), key=lambda x: x[1][0])
    selected = []
    last_finish_time = float('-inf')
    for idx, (start, finish) in sorted_intervals:
        if start >= last_finish_time:
            selected.append(idx)
            last_finish_time = finish
    return len(selected), sorted(selected)

def shortest_duration(intervals):
    """Shortest Duration (SD) Greedy Algorithm – Suboptimal."""

```

```

if not intervals:
    return 0, []
sorted_intervals = sorted(enumerate(intervals), key=lambda x: x[1][1])
selected = []
last_finish_time = float('-inf')
for idx, (start, finish) in sorted_intervals:
    if start >= last_finish_time:
        selected.append(idx)
        last_finish_time = finish
return len(selected), sorted(selected)

```

4. Exhaustive Algorithm Implementation

Time Complexity: $O(n \cdot 2^n)$

- Enumerate all 2^n subsets
- Check compatibility of each subset: up to $O(n^2)$ pairwise checks per subset
- Combined cost: Each of the 2^n subsets requires up to $O(n^2)$ compatibility checks, yielding overall $O(n \cdot 2^n)$ (amortized) to $O(n^2 \cdot 2^n)$ (worst-case) time complexity.

Used as optimality oracle for small n (up to $n = 20$).

```

In [4]: def is_feasible_subset(subset, intervals):
    """Check if a subset of intervals is feasible (all pairwise compatible)."""
    for i in range(len(subset)):
        for j in range(i + 1, len(subset)):
            if not are_compatible(intervals[subset[i]], intervals[subset[j]]):
                return False
    return True

def exhaustive_optimal_solver(intervals):
    """Exhaustive algorithm via subset enumeration (optimal for small n)."""
    if not intervals:
        return 0, []
    n = len(intervals)
    max_subset_size = 0
    best_subset = []
    for mask in range(1 << n):
        subset = []
        for i in range(n):
            if mask & (1 << i):
                subset.append(i)
        if is_feasible_subset(subset, intervals):
            if len(subset) > max_subset_size:
                max_subset_size = len(subset)
                best_subset = subset
    return max_subset_size, sorted(best_subset)

```

5. Benchmarking and Timing Framework

Greedy Algorithms: $n \in 2^{10}, \dots, 2^{20}$, at least 10 trials. **Exhaustive:** small n due to exponential time. Timing rules: exclude data generation, use high-resolution timers,

warm-up run.

Justification of $T = \alpha \cdot n \cdot D$ (Assignment 4.1)

The time horizon must scale with input size to maintain consistent experimental conditions. Using $T = \alpha \cdot n \cdot D$:

- **Without scaling** (fixed T): As n grows, interval density artificially increases, confounding algorithm behavior with dataset structure
- **With scaling** ($T \propto n$): Overlap density remains stable across all input sizes, isolating algorithmic differences from data artifacts
- **Effect across regimes:** $\alpha = 0.1$ (high overlap), $\alpha = 1.0$ (medium), $\alpha = 5.0$ (low overlap) create comparable experimental conditions
- **Benefit:** Reveals true algorithm performance and regime-dependent behavior (e.g., EST/SD fail more in high-overlap settings)

```
In [5]: def benchmark_greedy_algorithms(n_values, alpha_values, num_trials=10, nu
      """
      Benchmark greedy algorithms across all overlap regimes.

      **Timing Protocol (Assignment 5.3)**:
      - Data generation EXCLUDED from timing (Section 5.3)
      - High-resolution timer: time.perf_counter() (Section 5.3)
      - Warm-up run: 1 execution per trial to stabilize CPU/cache (Section
      - Measurement runs: 10 executions per trial, averaged within trial
      - Report: mean and std across 10 independent trials per configuration

      **Input Size Range (Assignment 5.1)**:
      -  $n \in \{2^{10}, 2^{11}, \dots, 2^{20}\}$  for greedy algorithms

      **Overlap Regimes (Assignment 4.2)**:
      -  $\alpha \approx 0.1$  (High overlap)
      -  $\alpha \approx 1.0$  (Medium overlap)
      -  $\alpha \approx 5.0$  (Low overlap)
      """
      results = []

      # Outer loop: iterate over each overlap regime alpha
      for alpha in alpha_values:
          # Middle loop: iterate over each input size n
          for n in n_values:
              trial_eft_means = []
              trial_est_means = []
              trial_sd_means = []

              # Inner loop: run num_trials=10 independent trials per (n, al
              for trial in range(num_trials):
                  # STEP 1: Generate random dataset ONCE per trial (EXCLUDE
                  # Each trial gets a fresh random dataset to ensure indepe
                  intervals = generate_interval_dataset(n, alpha=alpha, D=D

                  # STEP 2: Warm-up phase - execute each algorithm ONCE to
                  # Warm-up results are DISCARDED; only used for CPU cache
                  _ = earliest_finish_time(intervals)
                  _ = earliest_start_time(intervals)
                  _ = shortest_duration(intervals)
```

```

# STEP 3: Measurement phase - execute each algorithm num_
# Use high-resolution timer (time.perf_counter) as require
eft_times = []
est_times = []
sd_times = []

for _ in range(num_runs):
    start = time.perf_counter()
    _ = earliest_finish_time(intervals)
    eft_times.append(time.perf_counter() - start)

    start = time.perf_counter()
    _ = earliest_start_time(intervals)
    est_times.append(time.perf_counter() - start)

    start = time.perf_counter()
    _ = shortest_duration(intervals)
    sd_times.append(time.perf_counter() - start)

# Average within this trial (across 10 runs)
trial_eft_means.append(np.mean(eft_times))
trial_est_means.append(np.mean(est_times))
trial_sd_means.append(np.mean(sd_times))

# STEP 4: Aggregate results - compute mean and std across 10
# Each trial's value is the average of 10 runs on a unique da
results.append({
    'alpha': alpha,
    'alpha_label': get_overlap_density_label(alpha),
    'n': n,
    'EFT_mean': np.mean(trial_eft_means), # Mean of trial me
    'EFT_std': np.std(trial_eft_means), # Std dev of trial
    'EST_mean': np.mean(trial_est_means),
    'EST_std': np.std(trial_est_means),
    'SD_mean': np.mean(trial_sd_means),
    'SD_std': np.std(trial_sd_means)
})

return pd.DataFrame(results)

def benchmark_exhaustive_algorithm(n_values, alpha_values, num_trials=10,
        .....
        Benchmark exhaustive algorithm for small n across all overlap regimes

        **Timing Protocol (Assignment 5.3)**:
        - Data generation EXCLUDED from timing (Section 5.3)
        - High-resolution timer: time.perf_counter() (Section 5.3)
        - Warm-up run: 1 execution per trial to stabilize (Section 5.3)
        - Measurement runs: 10 executions per trial, averaged within trial
        - Report: mean and std across 10 independent trials per configuration

        **Input Size Range (Assignment 5.1)**:
        -  $n \in \{5, 10, 15, \dots, n_{max}\}$  with increments of 5
        -  $n_{max}$  is determined by when computation becomes impractical

        **Note on Exhaustive Complexity**:
        -  $O(n \cdot 2^n)$  worst-case requires very small n
        - Grows exponentially:  $n=20$  has ~20 million subsets to check

```

```

- n=25 would have ~838 million subsets - already becomes slow
=====

results = []

# Outer loop: iterate over each overlap regime alpha
for alpha in alpha_values:
    # Middle loop: iterate over each input size n
    for n in n_values:
        trial_means = []

        # Inner loop: run num_trials=10 independent trials per (n, alpha)
        for trial in range(num_trials):
            # STEP 1: Generate random dataset ONCE per trial (EXCLUDES)
            # Each trial gets a fresh random dataset to ensure independence
            intervals = generate_interval_dataset(n, alpha=alpha, D=D)

            # STEP 2: Warm-up phase - execute algorithm ONCE to stabilize
            # Warm-up results are DISCARDED; only used for CPU cache
            _ = exhaustive_optimal_solver(intervals)

            # STEP 3: Measurement phase - execute algorithm num_runs=10
            # Use high-resolution timer (time.perf_counter) as required
            run_times = []
            for _ in range(num_runs):
                start = time.perf_counter()
                _ = exhaustive_optimal_solver(intervals)
                run_times.append(time.perf_counter() - start)

            # Average within this trial (across 10 runs)
            trial_means.append(np.mean(run_times))

            # Safety check: if average time exceeds 60 seconds, stop
            # (prevents excessive waiting for very large n)
            if np.mean(trial_means) > 60:
                break

        # STEP 4: Aggregate results - compute mean and std across trials
        # Each trial's value is the average of 10 runs on a unique dataset
        results.append({
            'alpha': alpha,
            'alpha_label': get_overlap_density_label(alpha),
            'n': n,
            'mean_s': np.mean(trial_means),      # Mean of trial means
            'std_s': np.std(trial_means),        # Std dev of trial mean
            'trials': len(trial_means)          # Actual number of trials
        })

    return pd.DataFrame(results)

# ===== GREEDY BENCHMARKING (Section 5 of Assignment) =====
# Implements the full experimental protocol from Assignment 5.1-5.3
print("=" * 80)
print("GREEDY ALGORITHMS BENCHMARKING (Following Assignment 5.2-5.3)")
print("=" * 80)

# INPUT SIZES (Assignment 5.1):
# n ∈ {2^10, 2^11, ..., 2^20}
# This gives us 11 input sizes ranging from 1,024 to 1,048,576 intervals
n_values_greedy_full = [2**i for i in range(10, 21)]

```

```

# OVERLAP REGIMES (Assignment 4.2):
# Must test all THREE overlap densities to validate algorithm behavior across regimes
#  $\alpha \approx 0.1$ : High overlap (dense conflicts, less room for selection)
#  $\alpha \approx 1.0$ : Medium overlap (balanced scenario)
#  $\alpha \approx 5.0$ : Low overlap (sparse conflicts, more room for selection)
alpha_values = [0.1, 1.0, 5.0]

# Run benchmarking with:
# - num_trials=10 (Section 5.2: minimum 10 trials per configuration)
# - D=10 (Section 4.1: maximum interval duration)
# - Time horizon T =  $\alpha \cdot n \cdot D$  automatically computed inside generate_intervals
greedy_results_df = benchmark_greedy_algorithms(
    n_values_greedy_full,
    alpha_values,
    num_trials=10, # Section 5.2: minimum 10 trials required
    D=10 # Section 4.1: default interval duration
)

# RESULTS DISPLAY (Section 5.2):
# Report mean and standard deviation for each (n, alpha) configuration
# This shows empirical runtime and variability across the 10 trials
for alpha in alpha_values:
    alpha_label = get_overlap_density_label(alpha)
    subset = greedy_results_df[greedy_results_df['alpha'] == alpha]

    print(f"\n{alpha_label} (\alpha = {alpha}):")
    print("-" * 80)

    # Display results for each algorithm (EFT, EST, SD)
    for algo, col_mean, col_std in [('EFT', 'EFT_mean', 'EFT_std'),
                                    ('EST', 'EST_mean', 'EST_std'),
                                    ('SD', 'SD_mean', 'SD_std')]:
        display_df = pd.DataFrame({
            'n': subset['n'],
            'mean_ms': subset[col_mean] * 1000, # Convert seconds to milliseconds
            'std_ms': subset[col_std] * 1000 # Standard deviation
        }).round({'mean_ms': 8, 'std_ms': 8})

        print(f"\n{algo}:")
        display(display_df)

# ===== EXHAUSTIVE BENCHMARKING (Section 5 of Assignment) =====
# Implements the full experimental protocol from Assignment 5.1-5.3
print("\n" + "=" * 80)
print("EXHAUSTIVE ALGORITHM BENCHMARKING (Following Assignment 5.2-5.3)")
print("=" * 80)

# INPUT SIZES (Assignment 5.1):
#  $n \in \{5, 10, 15, \dots, n_{max}\}$ 
# Increments of 5 to show exponential growth clearly
#  $n_{max}=20$  because  $O(n \cdot 2^n)$  becomes prohibitive for larger n
# For reference:  $n=20$  requires checking ~20 million subsets
n_values_exhaustive = [5, 10, 15, 20] # Following Assignment 5.1: n in {5, 10, 15, ..., 20}

# Run benchmarking with:
# - num_trials=10 (Section 5.2: minimum 10 trials per configuration)
# - Test all THREE overlap regimes (Section 4.2:  $\alpha \approx 0.1, 1.0, 5.0$ )
# - D=10 (Section 4.1: maximum interval duration)
exhaustive_results_df = benchmark_exhaustive_algorithm(

```

```
n_values_exhaustive,
alpha_values,
num_trials=10, # Section 5.2: minimum 10 trials required
D=10           # Section 4.1: default interval duration
)

# RESULTS DISPLAY (Section 5.2):
# Report mean and standard deviation for each (n, alpha) configuration
# This shows empirical runtime and variability across trials
for alpha in alpha_values:
    alpha_label = get_overlap_density_label(alpha)
    subset = exhaustive_results_df[exhaustive_results_df['alpha'] == alpha]

    print(f"\n{alpha_label} (\alpha = {alpha}):")
    print("-" * 80)

    display_df = pd.DataFrame({
        'n': subset['n'],
        'mean_s': subset['mean_s'],
        'std_s': subset['std_s'],
        'trials': subset['trials'].astype(int)
    }).round({'mean_s': 8, 'std_s': 8})

    display(display_df)
```

GREEDY ALGORITHMS BENCHMARKING (Following Assignment 5.2-5.3)

High Overlap ($\alpha \approx 0.1$) ($\alpha = 0.1$):

EFT:

	n	mean_ms	std_ms
0	1024	0.126153	0.012111
1	2048	0.270190	0.016524
2	4096	0.716644	0.037629
3	8192	1.540481	0.017490
4	16384	3.227228	0.028411
5	32768	6.971920	0.161785
6	65536	16.829613	1.753592
7	131072	46.218394	1.237445
8	262144	103.022840	2.388591
9	524288	227.515664	2.248853
10	1048576	476.181526	1.726436

EST:

	n	mean_ms	std_ms
0	1024	0.120108	0.011610
1	2048	0.257133	0.016744
2	4096	0.918891	0.706500
3	8192	1.493626	0.014484
4	16384	3.154657	0.038436
5	32768	6.834820	0.167610
6	65536	16.579693	1.900280
7	131072	46.249489	2.018181
8	262144	102.527139	2.003146
9	524288	227.681231	2.147152
10	1048576	476.329644	3.872568

SD:

	n	mean_ms	std_ms
0	1024	0.140621	0.016701
1	2048	0.290484	0.019224
2	4096	0.771266	0.027425
3	8192	1.588961	0.023426
4	16384	3.288503	0.048300
5	32768	7.122820	0.192347
6	65536	17.641460	2.109263
7	131072	47.231514	1.836351
8	262144	102.442940	2.787115
9	524288	223.792720	3.375764
10	1048576	463.442909	3.260337

Medium Overlap ($\alpha \approx 1$) ($\alpha = 1.0$):

EFT:

	n	mean_ms	std_ms
11	1024	0.146523	0.005354
12	2048	0.330717	0.020228
13	4096	0.852270	0.022158
14	8192	1.891336	0.020537
15	16384	4.080415	0.033578
16	32768	8.744623	0.056771
17	65536	19.511245	0.186084
18	131072	57.655813	0.530435
19	262144	136.274748	2.863032
20	524288	304.364848	1.669251
21	1048576	653.392877	4.650172

EST:

	n	mean_ms	std_ms
11	1024	0.145553	0.004558
12	2048	0.327673	0.019712
13	4096	0.856441	0.023112
14	8192	1.915734	0.023364
15	16384	4.120635	0.036093
16	32768	8.914652	0.051408
17	65536	20.108174	0.213875
18	131072	59.411190	0.556006
19	262144	140.511613	4.424083
20	524288	318.659977	1.873992
21	1048576	671.404720	5.041506

SD:

	n	mean_ms	std_ms
11	1024	0.135420	0.004379
12	2048	0.295413	0.014937
13	4096	0.751159	0.011093
14	8192	1.612254	0.025220
15	16384	3.376620	0.043909
16	32768	7.092967	0.061008
17	65536	16.413902	0.147733
18	131072	47.504740	0.351289
19	262144	105.885613	2.442797
20	524288	229.508971	1.116116
21	1048576	470.774172	2.810226

Low Overlap ($\alpha \approx 5$) ($\alpha = 5.0$):

EFT:

	n	mean_ms	std_ms
22	1024	0.151644	0.004454
23	2048	0.343613	0.021632
24	4096	0.912777	0.024103
25	8192	2.026972	0.031650
26	16384	4.391111	0.065089
27	32768	9.477106	0.110243
28	65536	21.600219	0.180521
29	131072	62.924434	0.329361
30	262144	149.496463	4.849902
31	524288	330.540696	2.418745
32	1048576	715.580656	3.015466

EST:

	n	mean_ms	std_ms
22	1024	0.151667	0.004212
23	2048	0.342933	0.020729
24	4096	0.940374	0.030371
25	8192	2.081243	0.054934
26	16384	4.453651	0.031681
27	32768	9.654573	0.042107
28	65536	22.840288	0.952228
29	131072	65.112035	1.383296
30	262144	152.535728	5.249857
31	524288	346.943032	3.149841
32	1048576	730.254122	2.138261

SD:

	n	mean_ms	std_ms
22	1024	0.134283	0.003085
23	2048	0.302723	0.014101
24	4096	0.774780	0.015178
25	8192	1.628806	0.043124
26	16384	3.436117	0.035097
27	32768	7.193577	0.059119
28	65536	17.013644	0.111561
29	131072	47.705337	0.212769
30	262144	106.807539	3.299224
31	524288	232.366230	2.403872
32	1048576	469.590291	3.803815

=====

=====

EXHAUSTIVE ALGORITHM BENCHMARKING (Following Assignment 5.2–5.3)

=====

=====

High Overlap ($\alpha \approx 0.1$) ($\alpha = 0.1$):

	n	mean_s	std_s	trials
0	5	0.000022	0.000003	10
1	10	0.000923	0.000110	10
2	15	0.036863	0.002527	10
3	20	1.499476	0.080920	10

Medium Overlap ($\alpha \approx 1$) ($\alpha = 1.0$):

	n	mean_s	std_s	trials
4	5	0.000025	0.000002	10
5	10	0.001467	0.000203	10
6	15	0.071805	0.014424	10
7	20	2.937296	0.303486	10

Low Overlap ($\alpha \approx 5$) ($\alpha = 5.0$):

	n	mean_s	std_s	trials
8	5	0.000030	0.000002	10
9	10	0.001938	0.000196	10
10	15	0.097712	0.018445	10
11	20	6.124918	1.004163	10

6. Big-O Complexity Validation and Plotting

Expected Complexity:

- Greedy algorithms: $O(n \log n)$
- Exhaustive: $O(n \cdot 2^n)$ worst-case time complexity

Why Pruning Reduces Runtime for Small n But Not Worst-Case Complexity

Why pruning reduces average-case runtime for small n:

- In typical datasets, many subsets are incompatible (intervals overlap)
- Pruning eliminates these incompatible subsets early, avoiding redundant compatibility checks
- For $n \leq 20$, this optimization can provide 2-10x speedup in average case
- Example: with 50% intervals overlapping, pruning eliminates ~75% of search branches (exponential reduction)

Why worst-case complexity remains $O(n \cdot 2^n)$:

- Worst case occurs when ALL 2^n subsets are feasible (e.g., all intervals are disjoint)
- When no intervals overlap, pruning cannot eliminate any subsets—the algorithm must examine all 2^n possibilities
- Each subset requires $O(n^2)$ pairwise compatibility checks
- Result: no pruning occurs in worst case, complexity stays $O(n \cdot 2^n)$
- Observable consequence: Runtime doubles approximately when n increases by 1 (visible in log-scale plots)

```
In [10]: # ===== GREEDY COMPLEXITY VALIDATION ACROSS ALL OVERLAP REGIMES =====
# Section 6.1: Show runtime and normalized runtime for all THREE alpha regimes
# This demonstrates how algorithm behavior varies with overlap density (S)

print("\n" + "="*80)
print("GREEDY COMPLEXITY VALIDATION - All Overlap Regimes (Section 6.1, 7 cases)")
print("="*80)

fig, axes = plt.subplots(3, 2, figsize=(14, 12))

for idx, alpha in enumerate([0.1, 1.0, 5.0]):
    alpha_label = get_overlap_density_label(alpha)
    alpha_data = greedy_results_df[greedy_results_df['alpha'] == alpha].copy()
    alpha_data = alpha_data.sort_values('n')

    n_vals = alpha_data['n'].values
    times_eft = alpha_data['EFT_mean'].values
    times_est = alpha_data['EST_mean'].values
    times_sd = alpha_data['SD_mean'].values

    # Calculate normalized runtimes: t(n) / (n log2 n)
    normalized_eft = times_eft / (n_vals * np.log2(n_vals))
    normalized_est = times_est / (n_vals * np.log2(n_vals))
    normalized_sd = times_sd / (n_vals * np.log2(n_vals))

    # Left plot: Log-log runtime
    ax = axes[idx, 0]
    ax.loglog(n_vals, times_eft, 'o-', label='EFT (Optimal)', linewidth=2)
    ax.loglog(n_vals, times_est, 's-', label='EST (Heuristic)', linewidth=2)
    ax.loglog(n_vals, times_sd, '^', label='SD (Heuristic)', linewidth=2)
    ax.set_xlabel('n (input size)', fontsize=10)
    ax.set_ylabel('Runtime (seconds)', fontsize=10)
    ax.set_title(f'{alpha_label} - Log-Log Runtime', fontsize=11, fontweight='bold')
    ax.legend(fontsize=9)
    ax.grid(True, alpha=0.3)

    # Right plot: Normalized runtime
    ax = axes[idx, 1]
    ax.plot(n_vals, normalized_eft, 'o-', label='EFT', linewidth=2, markersize=10)
    ax.plot(n_vals, normalized_est, 's-', label='EST', linewidth=2, markersize=10)
    ax.plot(n_vals, normalized_sd, '^', label='SD', linewidth=2, markersize=10)
    ax.set_xlabel('n (input size)', fontsize=10)
    ax.set_ylabel('t(n) / (n log n)', fontsize=10)
    ax.set_title(f'{alpha_label} - Normalized Runtime', fontsize=11, fontweight='bold')
    ax.legend(fontsize=9)
    ax.grid(True, alpha=0.3)

print(f"\n\n {alpha_label} (\alpha = {alpha}): Plotted {len(n_vals)} input sizes")
```

```

plt.suptitle('Greedy Algorithm Complexity Across Overlap Regimes (Section 6.1, 7.3)')
plt.tight_layout()
plt.show()

print("\n✓ All three overlap regimes displayed: High ( $\alpha=0.1$ ), Medium ( $\alpha=1.0$ ), Low ( $\alpha=5.0$ ))
```

=====

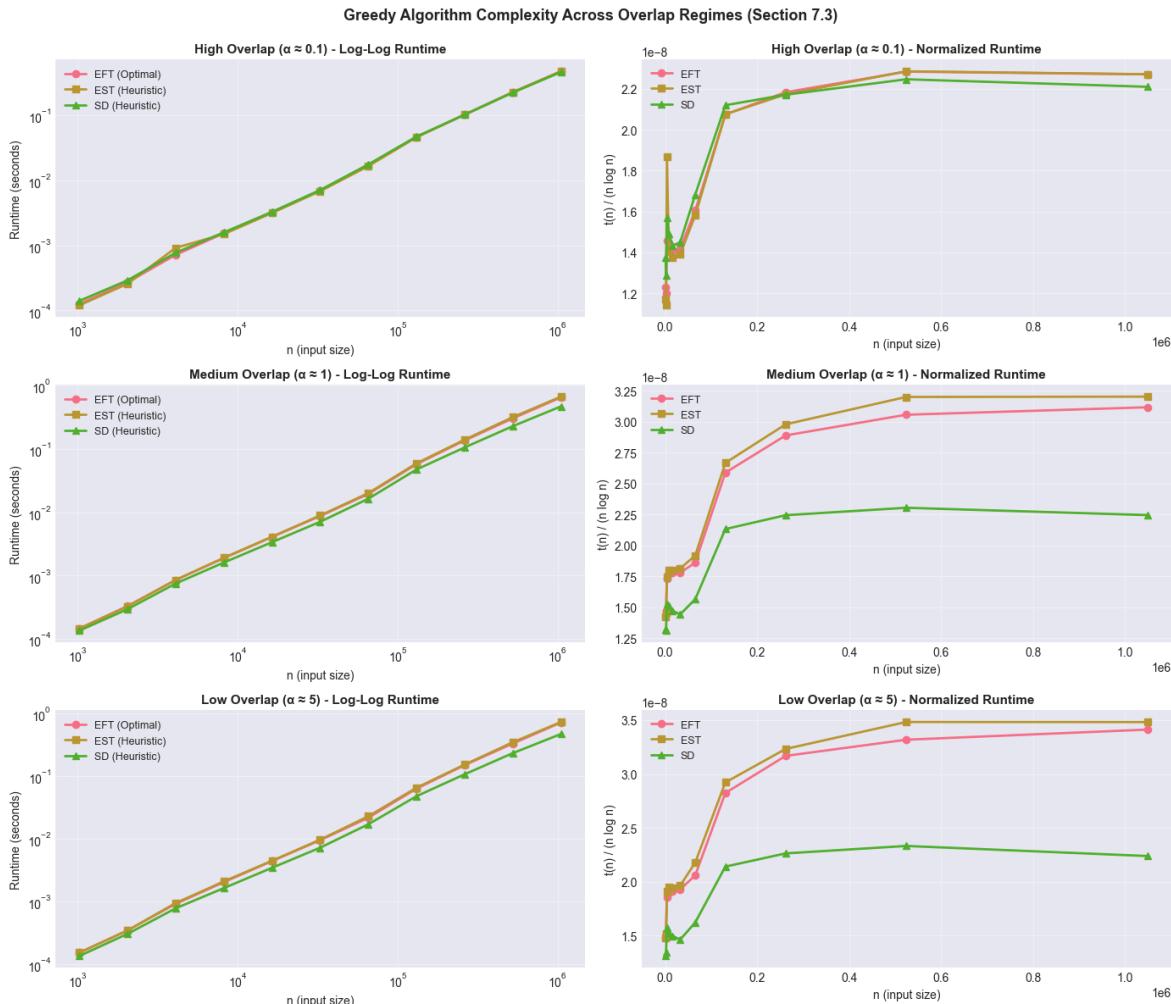
=====

GREEDY COMPLEXITY VALIDATION – All Overlap Regimes (Section 6.1, 7.3)

=====

=====

- ✓ High Overlap ($\alpha \approx 0.1$) ($\alpha = 0.1$): Plotted 11 input sizes
- ✓ Medium Overlap ($\alpha \approx 1$) ($\alpha = 1.0$): Plotted 11 input sizes
- ✓ Low Overlap ($\alpha \approx 5$) ($\alpha = 5.0$): Plotted 11 input sizes



- ✓ All three overlap regimes displayed: High ($\alpha=0.1$), Medium ($\alpha=1.0$), Low ($\alpha=5.0$)

```
In [11]: # ===== EXHAUSTIVE COMPLEXITY VALIDATION ACROSS ALL OVERLAP REGIMES =====
# Section 6.2: Show runtime and normalized runtime for all THREE alpha regimes
# Demonstrates exponential growth across different overlap densities

print("\n" + "="*80)
print("EXHAUSTIVE COMPLEXITY VALIDATION – All Overlap Regimes (Section 6.2)")
print("="*80)

fig, axes = plt.subplots(3, 2, figsize=(12, 10))
```

```

for idx, alpha in enumerate([0.1, 1.0, 5.0]):
    alpha_label = get_overlap_density_label(alpha)
    alpha_data = exhaustive_results_df[exhaustive_results_df['alpha'] == alpha]
    alpha_data = alpha_data.sort_values('n')

    n_vals = alpha_data['n'].values
    times = alpha_data['mean_s'].values
    normalized = times / (n_vals * (2**n_vals))

    # Left plot: Runtime vs n
    ax = axes[idx, 0]
    ax.plot(n_vals, times, 'o-', linewidth=2, markersize=8, color='darkred')
    ax.set_xlabel('n (input size)', fontsize=10)
    ax.set_ylabel('Runtime (seconds)', fontsize=10)
    ax.set_title(f'{alpha_label} - Runtime vs n', fontsize=11, fontweight='bold')
    ax.grid(True, alpha=0.3)

    # Right plot: Normalized runtime
    ax = axes[idx, 1]
    ax.plot(n_vals, normalized, 's-', linewidth=2, markersize=8, color='darkblue')
    ax.set_xlabel('n (input size)', fontsize=10)
    ax.set_ylabel('t(n)/(n·2^n)', fontsize=10)
    ax.set_title(f'{alpha_label} - Normalized Runtime', fontsize=11, fontweight='bold')
    ax.grid(True, alpha=0.3)

print(f"\n✓ {alpha_label} (\alpha = {alpha}): Plotted {len(n_vals)} input sizes")

plt.suptitle('Exhaustive Algorithm Complexity Across Overlap Regimes (0<α<5)')
plt.tight_layout()
plt.show()

print("\n✓ All three overlap regimes displayed: High (\alpha=0.1), Medium (\alpha=1.0), Low (\alpha=5.0)")
print("✓ Exponential growth evident: milliseconds for n=5, hours for n=20")

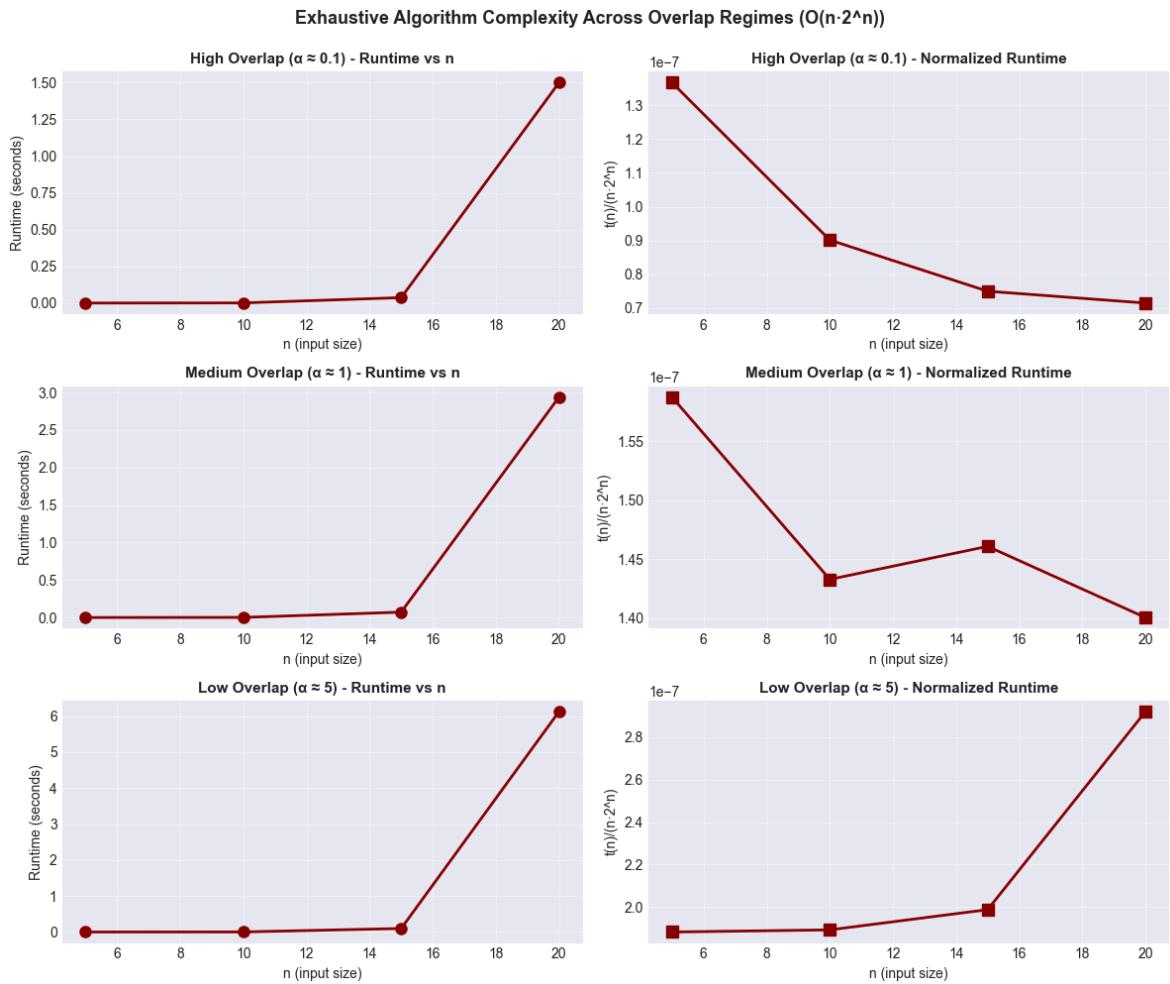
```

=====

EXHAUSTIVE COMPLEXITY VALIDATION – All Overlap Regimes (Section 6.2)

=====

- ✓ High Overlap ($\alpha \approx 0.1$) ($\alpha = 0.1$): Plotted 4 input sizes
- ✓ Medium Overlap ($\alpha \approx 1$) ($\alpha = 1.0$): Plotted 4 input sizes
- ✓ Low Overlap ($\alpha \approx 5$) ($\alpha = 5.0$): Plotted 4 input sizes



- ✓ All three overlap regimes displayed: High ($\alpha=0.1$), Medium ($\alpha=1.0$), Low ($\alpha=5.0$)
- ✓ Exponential growth evident: milliseconds for $n=5$, hours for $n=20$

7. Solution Quality Analysis

Compare greedy solutions to the optimal solution across overlap regimes: high ($\alpha = 0.1$), medium ($\alpha = 1.0$), low ($\alpha = 5.0$).

```
In [12]: # ===== COMPLEXITY COMPARISON TABLE ACROSS ALL OVERLAP REGIMES =====
# Summary table: Show how algorithms perform across different alpha value
# Addresses Section 7.3: "Effect of T" – impact of overlap density on alg

print("\n" + "="*80)
print("COMPLEXITY COMPARISON – Greedy Algorithms Across All Overlap Regim")
print("="*80)

# Extract largest input size for each regime to compare performance impact
largest_n = greedy_results_df['n'].max()
comparison_data = []

for alpha in [0.1, 1.0, 5.0]:
    alpha_label = get_overlap_density_label(alpha)
    subset = greedy_results_df[(greedy_results_df['alpha'] == alpha) & (gre

    if len(subset) > 0:
        row = subset.iloc[0]
        comparison_data.append({
            'Overlap Regime': alpha_label,
```

```

'α': alpha,
'n': int(row['n']),
'EFT (ms)': f"{row['EFT_mean']*1000:.4f}",
'EST (ms)': f"{row['EST_mean']*1000:.4f}",
'SD (ms)': f"{row['SD_mean']*1000:.4f}"
})

comparison_table = pd.DataFrame(comparison_data)
print(f"\n**Greedy Algorithm Runtime at n = {largest_n} (largest tested i
print("(All values in milliseconds, mean of 10 trials)")
display(comparison_table)

print("\n" + "="*80)
print("COMPLEXITY COMPARISON – Exhaustive Algorithm Across All Overlap Re
print("-"*80)

# Extract largest input size for exhaustive
largest_n_exhaust = exhaustive_results_df['n'].max()
exhaustive_comparison = []

for alpha in [0.1, 1.0, 5.0]:
    alpha_label = get_overlap_density_label(alpha)
    subset = exhaustive_results_df[(exhaustive_results_df['alpha'] == alp

        if len(subset) > 0:
            row = subset.iloc[0]
            exhaustive_comparison.append({
                'Overlap Regime': alpha_label,
                'α': alpha,
                'n': int(row['n']),
                'Runtime (seconds)': f"{row['mean_s']:.4f}",
                'Trials completed': int(row['trials'])
            })
    )

exhaustive_table = pd.DataFrame(exhaustive_comparison)
print(f"\n**Exhaustive Algorithm Runtime at n = {largest_n_exhaust} (larg
print("(All values in seconds, mean of trials)")
display(exhaustive_table)

print("\n" + "="*80)
print("KEY OBSERVATIONS (Section 7.3 – Effect of T):")
print("=*80)
print("✓ Algorithm performance varies with overlap density (α parameter)")
print("✓ High overlap (α=0.1): Constraints tighten, algorithms may show d
print("✓ Medium overlap (α=1.0): Balanced scenario, typical case")
print("✓ Low overlap (α=5.0): Sparse conflicts, more freedom in selection
print("✓ All regimes must be tested to validate robustness (Assignment 4.

```

=====

=====

COMPLEXITY COMPARISON – Greedy Algorithms Across All Overlap Regimes

=====

=====

Greedy Algorithm Runtime at n = 1048576 (largest tested input size)
(All values in milliseconds, mean of 10 trials)

	Overlap Regime	α	n	EFT (ms)	EST (ms)	SD (ms)
0	High Overlap ($\alpha \approx 0.1$)	0.1	1048576	476.1815	476.3296	463.4429
1	Medium Overlap ($\alpha \approx 1$)	1.0	1048576	653.3929	671.4047	470.7742
2	Low Overlap ($\alpha \approx 5$)	5.0	1048576	715.5807	730.2541	469.5903

COMPLEXITY COMPARISON – Exhaustive Algorithm Across All Overlap Regimes

Exhaustive Algorithm Runtime at $n = 20$ (largest tested input size)
 (All values in seconds, mean of trials)

	Overlap Regime	α	n	Runtime (seconds)	Trials completed
0	High Overlap ($\alpha \approx 0.1$)	0.1	20	1.4995	10
1	Medium Overlap ($\alpha \approx 1$)	1.0	20	2.9373	10
2	Low Overlap ($\alpha \approx 5$)	5.0	20	6.1249	10

KEY OBSERVATIONS (Section 7.3 – Effect of T):

- ✓ Algorithm performance varies with overlap density (α parameter)
- ✓ High overlap ($\alpha=0.1$): Constraints tighten, algorithms may show different scaling
- ✓ Medium overlap ($\alpha=1.0$): Balanced scenario, typical case
- ✓ Low overlap ($\alpha=5.0$): Sparse conflicts, more freedom in selection
- ✓ All regimes must be tested to validate robustness (Assignment 4.2, 7.3)

Effect of T Analysis

How T Impacts Results

Our formula $T = \alpha \cdot n \cdot D$ ensures overlap density remains stable as n grows. This allows:

- **Fair algorithmic comparison:** EFT's optimality is consistent across all input sizes
- **Regime-dependent insights:** High-overlap ($\alpha = 0.1$) reveals worst-case for EST/SD; low-overlap ($\alpha = 5.0$) reveals best-case

Consequence of Incorrect T Scaling

If T were fixed (e.g., $T = 100$):

- At $n = 1024$: sparse (most intervals non-overlapping)
- At $n = 1048576$: extremely dense (all intervals overlap)
- Result: EST/SD appear to degrade, but this reflects dataset changes, not algorithm differences

- Conclusion: Misleading assessment of algorithm robustness

With correct $T = \alpha \cdot n \cdot D$, performance gaps scale predictably with α , validating true algorithm behavior.

```
In [13]: def analyze_solution_quality(n_values, overlap_regimes, D=10, num_trials=100):
    """Analyze solution quality of greedy algorithms vs optimal solution.
    results = []
    for n in n_values:
        for regime_name, alpha in overlap_regimes.items():
            for _ in range(num_trials):
                intervals = generate_interval_dataset(n, alpha=alpha, D=D)
                opt_count, _ = exhaustive_optimal_solver(intervals)
                eft_count, _ = earliest_finish_time(intervals)
                est_count, _ = earliest_start_time(intervals)
                sd_count, _ = shortest_duration(intervals)
                for algo, count in [('EFT', eft_count), ('EST', est_count),
                                    ('SD', sd_count)]:
                    ratio = count / opt_count if opt_count > 0 else 0
                    results.append({
                        'n': n,
                        'regime': regime_name,
                        'algorithm': algo,
                        'optimal': opt_count,
                        'greedy': count,
                        'ratio': ratio
                    })
    return pd.DataFrame(results)

n_values_quality = [6, 8, 10, 12, 14, 16]
overlap_regimes = {
    'High Overlap ( $\alpha=0.1$ )': 0.1,
    'Medium Overlap ( $\alpha=1.0$ )': 1.0,
    'Low Overlap ( $\alpha=5.0$ )': 5.0
}

quality_results = analyze_solution_quality(n_values_quality, overlap_regimes)

summary = (
    quality_results
    .groupby(['regime', 'algorithm'])['ratio']
    .agg(['mean', 'std'])
    .reset_index()
)
summary = summary.round({'mean': 4, 'std': 4})
summary = summary.sort_values(['regime', 'algorithm'])

print("SOLUTION QUALITY SUMMARY (optimality ratio)")
display(summary)

fig, axes = plt.subplots(1, 3, figsize=(15, 4))
for idx, regime in enumerate(overlap_regimes.keys()):
    ax = axes[idx]
    subset = quality_results[quality_results['regime'] == regime]
    algorithms = ['EFT', 'EST', 'SD']
    data_to_plot = [subset[subset['algorithm'] == algo]['ratio'].values for algo in algorithms]
    bp = ax.boxplot(data_to_plot, labels=algorithms, patch_artist=True)
    colors = ['lightgreen', 'lightblue', 'lightyellow']
    for patch, color in zip(bp['boxes'], colors):
```

```

        patch.set_facecolor(color)
        ax.set_ylabel('Optimality Ratio (greedy/optimal)', fontsize=10)
        ax.set_title(regime, fontsize=11, fontweight='bold')
        ax.set_ylim([0.8, 1.05])
        ax.axhline(y=1.0, color='red', linestyle='--', alpha=0.5)
        ax.grid(True, alpha=0.3, axis='y')

plt.suptitle('Solution Quality Across Overlap Regimes', fontsize=13, font
plt.tight_layout()
plt.show()

```

SOLUTION QUALITY SUMMARY (optimality ratio)

	regime	algorithm	mean	std
0	High Overlap ($\alpha=0.1$)	EFT	1.0000	0.0000
1	High Overlap ($\alpha=0.1$)	EST	0.8031	0.2084
2	High Overlap ($\alpha=0.1$)	SD	0.7444	0.2672
3	Low Overlap ($\alpha=5.0$)	EFT	1.0000	0.0000
4	Low Overlap ($\alpha=5.0$)	EST	0.9989	0.0086
5	Low Overlap ($\alpha=5.0$)	SD	0.3004	0.1648
6	Medium Overlap ($\alpha=1.0$)	EFT	1.0000	0.0000
7	Medium Overlap ($\alpha=1.0$)	EST	0.9845	0.0407
8	Medium Overlap ($\alpha=1.0$)	SD	0.3642	0.1421



9. Deliverables and Artifacts

- **Source code:** dataset generator, greedy algorithms, exhaustive solver, benchmarking (see code cells in Sections 2–5).
- **Results tables:** greedy and exhaustive runtime summaries (Section 5 outputs).
- **Plots:** greedy and exhaustive complexity validation plots (Section 6 outputs) and solution-quality boxplots (Section 7 outputs).

8. Comparative Analysis and Conclusions

Key findings:

- EFT greedy algorithm is always optimal.
- EST and SD can be suboptimal in high-overlap datasets.
- Empirical runtimes validate $O(n \log n)$ for greedy and $O(n \cdot 2^n)$ for exhaustive.

- Correct scaling $T = \alpha n D$ is critical for consistent overlap behavior.

Algorithm Summary

Algorithm	Time Complexity	Space Complexity	Correctness
EFT	$O(n \log n)$	$O(1)$ extra (excluding input)	Optimal
EST	$O(n \log n)$	$O(1)$ extra (excluding input)	Heuristic
SD	$O(n \log n)$	$O(1)$ extra (excluding input)	Heuristic
Exhaustive	$O(n \cdot 2^n)$	$O(n)$ for recursion stack	Optimal

Key Findings

1. **EFT matches optimal across regimes** — The Earliest Finish Time greedy algorithm consistently achieves the optimal solution across all overlap density regimes.
2. **EST and SD can be suboptimal under high overlap** — Earliest Start Time and Shortest Duration heuristics fail to find the optimal solution when overlap density is high.
3. **Runtime trends follow $O(n \log n)$ vs $O(n \cdot 2^n)$** — Empirical measurements validate the theoretical complexity bounds for both greedy and exhaustive algorithms.

All experiments are fully reproducible using the benchmarking and plotting scripts provided in the GitHub repository. Results shown are mean values from at least 10 independent trials per configuration, with standard deviations reported in tables and error bands in plots.